

Cartwheel Homework 6

Repeatable agent evaluations in Harbor and GitHub Actions

11 classified cases	55 baseline runs	110 CI agent runs	15 capability analysis runs
---	--	---	---

Outcome

We converted two observed Homework 4 failure modes into 11 trace-level cases, established five-run baselines, implemented pass@k and pass^k, and ran two full CI evaluations from one pull request. The intentional email-change regression failed 0/5 and recovered to 5/5 after its prompt instruction was removed.

Model: **gpt-5.6-luna** - Service tier: **Fast** - Judge calls: **0** - Final local commit: **0a9d6a1** - Report date: **September 26, 2026**

This screenshot report replaces the optional walkthrough video for the student's course record. The committed cases, Harbor result summaries, CI URLs, and capability analysis remain the authoritative evidence.

1. Cases and five-run baselines

Failure mode 1: missing escalation

Cases e-001 through e-005 cover account changes, charge disputes, and payment credential changes. Case e-011 probes an invalid negative product price that should trigger human escalation.

Failure mode 2: missing policy grounding

Cases e-006 through e-010 require the correct help-center lookup and exact policy identifier. All decisions use deterministic tool-call and reply-content checks.

Cases	Mode	Baseline	Classification
e-001 to e-005	Missing required escalation	5/5 each	Regression
e-006 to e-010	Missing policy grounding	5/5 each	Regression
e-011	Missing required escalation	0/5	Capability

Why e-011 is a capability

All five baseline replies noticed that the seeded -\$5.00 price was invalid, but none called `escalate_to_human`. The desired behavior was therefore not reliable enough to protect as a regression.

Scope and cost controls

The final set contains 11 cases, 55 baseline agent runs, and no LLM judge calls. Exact facts are checked in code, and the rejected Homework 5 judge is not used for autonomous enforcement.

2. Reliability metrics and CI contract

Component	Behavior
pass@k	Estimates the chance of at least one success in k attempts from n observed runs with c successes.
pass^k	Estimates the chance that all k attempts succeed. This is the reliability metric used for regression cases.
case_passes	Blocks CI when any completed regression trial fails. Capability cases report metrics without blocking CI.
Harbor adapter	Reads Harbor 0.23 per-trial result files while retaining compatibility with the earlier aggregate layout.

GitHub Actions

Every pull request installs Harbor 0.23, generates the classified tasks, runs five Docker-isolated attempts per case, summarizes even after failures, and uploads the full Harbor job artifact. CARTWHEEL_MODEL is a repository variable; provider credentials remain encrypted secrets.

Validation

The completed branch passed 86 offline checks locally. The pass metric and Harbor adapter subset passed 40 checks. A timing-sensitive assertion on an unrelated SDK warning was removed while preserving the direct no-export assertion; the final GitHub offline workflow passed.

Final free validation: github.com/SCRKT/cartwheel-homeworks/actions/runs/36236561526. The Harbor evaluation job was skipped because the demonstration pull request had been closed after its two required paid runs.

3. One pull request, two evaluation commits

The public fork records the complete HW6 evaluation pipeline. Commit 7942fc6 adds the intentional e-001 regression; commit 85084d2 removes it on the same pull request. The PR was closed after both required paid runs to prevent later artifact-only pushes from starting a third batch.

Platform ▾Solutions ▾Resources ▾Open Source ▾Enterprise ▾Pricing

Search /

Sign in

Sign up

⌵

SCRKT / cartwheel-homeworks Public

Notifications

Fork 0

Star 0

<> Code

Pull requests

Actions

Projects

Security and quality

Insights

Complete HW6 Harbor CI evaluation #1

🔒 Closed

SCRKT wants to merge 7 commits into `main` from `hw6-ci-clean`

💬 Conversation 0

📄 Commits 7

🔍 Checks 4

📄 Files changed 8

+298 -26

SCRKT commented 7 hours ago

Owner

...

Adds the Homework 6 evaluation pipeline and turns reviewed Homework 4 failures into repeatable Harbor cases.

- Adds 11 cases across missing escalation and missing policy grounding, classified from five baseline attempts each.
- Implements `pass@k`, `pass^k`, and regression-versus-capability CI decisions.
- Supports Harbor 0.23's per-trial result layout and adds regression coverage.
- Runs five attempts per case in GitHub Actions with `gpt-5.6-luna` on the Fast service tier.
- Includes the temporary `e-001` email-change regression required for the first Part D run. A follow-up commit on this PR will revert it for the second run.

Validation: 86 passed, 14 skipped, 13 xfailed, 3 xpassed .

👤 Samuel Crockett added 7 commits 8 hours ago

🔍 Draft HW6 baseline evaluation cases

61137f1

Development

🔍 Support Harbor 0.23 per-trial results

b8bec29

Successfully merging this pull request may close these issues.

🔍 Classify initial HW6 cases and add capability probe

fe9dc7f

None yet

🔍 Complete HW6 baselines and pass metrics

f14a286

1 participant

🔍 Configure HW6 Harbor evaluation workflow

d38a22c

🔍 Introduce HW6 email escalation regression

✖ 7942fc6

🔍 Revert intentional HW6 email regression

✖ 85084d2

Cartwheel - Homework 6 CI for agent evaluations

Page 4

4. Intentional regression run

The first pull-request run completed all 55 Harbor trials and uploaded one job artifact. The workflow failed at the summary step because e-001 scored 0/5 and was correctly classified as a blocking regression. Offline checks passed. Offline checks passed.

Platform Solutions Resources Open Source Enterprise Pricing

Search / Sign in Sign up

SCRKT / cartwheel-homeworks Public

forked from ai-evals-course/cartwheel-homeworks

Notifications Fork 0 Star 0

<> Code

Pull requests

Actions

Projects

Security and quality

Insights

← evals

Complete HW6 Harbor CI evaluation #2

Sign in to view logs

Summary

All jobs

offline checks

Harbor evaluations

Run details

Usage

Workflow file

Triggered via pull request 7 hours ago

Status

Total duration

Artifacts

SCRKT opened #1 hw6-ci-clean

Failure

16m 32s

1

evals.yml

on: pull_request

offline checks 21s

Harbor evaluations 16m 28s

Annotations

1 error, 2 warnings, and 2 notices

Harbor evaluations

Process completed with exit code 1.

offline checks

Node.js 20 is deprecated. The following actions target Node.js 20 but are being forced to run on Node.js 24: actions/checkout@v4, astral-sh/setup-uv@v5. For m...

Show more

Harbor evaluations

Node.js 20 is deprecated. The following actions target Node.js 20 but are being forced to run on Node.js 24: actions/checkout@v4, actions/upload-artifact@v4, ...

Show more

5. Reverted regression run

The second run used the corrected prompt. e-001 recovered to 5/5. A trace audit found one real e-003 behavioral failure and an e-006 evaluator false negative: the literal check for delivery rejected the equivalent word delivered. The separate offline failure was a timing-sensitive logging assertion fixed and validated afterward.

Platform Solutions Resources Open Source Enterprise Pricing

Search / Sign in Sign up

SCRKT / cartwheel-homeworks Public

forked from ai-evals-course/cartwheel-homeworks

Notifications Fork 0 Star 0

<> Code Pull requests Actions Projects Security and quality Insights

← evals

Complete HW6 Harbor CI evaluation #4 Sign in to view logs

Summary

All jobs

offline checks Harbor evaluations

Run details Usage Workflow file

Triggered via pull request 6 hours ago

Status

Total duration

Artifacts

SCRKT synchronize #1 hw6-ci-clean Failure 13m 12s 1

evals.yml

on: pull_request

offline checks 24s

Harbor evaluations 13m 7s

Annotations

2 errors, 2 warnings, and 2 notices

offline checks Process completed with exit code 1.

Harbor evaluations Process completed with exit code 1.

offline checks Node.js 20 is deprecated. The following actions target Node.js 20 but are being forced to run on Node.js 24: actions/checkout@v4, astral-sh/setup-uv@v5. For m... Show more

Harbor evaluations

6. CI results and interpretation

Case	First run	After revert	Interpretation
e-001	0/5 - block	5/5 - pass	Intentional regression recovered
e-002	5/5	5/5	Reliable in both runs
e-003	5/5	4/5 - block	New reliability failure
e-004	5/5	5/5	Reliable in both runs
e-005	5/5	5/5	Reliable in both runs
e-006	3/5 - block	4/5 - block	False negatives from delivery/delivered check
e-007 to e-010	5/5 each	5/5 each	Reliable in both runs
e-011 capability	2/5	1/5	Reported, never blocks

Run 1

Commit 7942fc6 - 55 agent runs - 0 judge calls - 16m32s - artifact uploaded. URL: github.com/SCRKT/cartwheel-homeworks/actions/runs/36215130310

Run 2

Commit 85084d2 - 55 agent runs - 0 judge calls - 13m12s - artifact uploaded. URL: github.com/SCRKT/cartwheel-homeworks/actions/runs/36216043198

Post-run audit: Re-evaluating all 15 saved e-006 transcripts with the shared stem **deliver** produces 15/15 passes. Historical CI results remain unchanged; the corrected interpretation treats those three failures as evaluator defects.

7. Capability case after 5, 10, and 15 runs

Observed runs	Successes	pass@1	pass@3	pass@5	pass@10	pass@15
5	2	0.400	0.900	1.000	-	-
10	4	0.400	0.833	0.976	-	-
15	4	0.267	0.637	0.846	0.996	1.000

Stability conclusion

The final five trials all failed, moving pass@1 from 0.400 to 0.267, pass@3 from 0.833 to 0.637, and pass@5 from 0.976 to 0.846. These changes are material, so 15 observed runs did not produce a stable estimate. A product should not rely on repeated attempts to compensate for the missing escalation behavior.

Artifact	Location or result
Cases	eval_cases/cases.jsonl - 11 classified cases across two failure modes
Baselines	eval_results/hw6-baseline.json - 55 completed trials, zero exceptions
CI evidence	ci-runs.json - both run URLs, commits, targeted outcomes, and explanation
Capability analysis	eval_results/e-011-15.json - ordered trials and 5/10/15 estimates
Workflow	.github/workflows/evals.yml - Harbor 0.23, Docker, five attempts, summary and artifact
Fork branch	SCRKT/cartwheel-homeworks:hw6-ci-clean - final commit ed01b9b